

Unit-III: Neural Networks

Unit-III: Neural Networks: Anatomy of Neural Network, Introduction to Keras: Keras, TensorFlow, Theano and CNTK, setting up Deep Learning Workstation, Classifying Movie Reviews: Binary Classification, Classifying newswires: Multiclass Classification.

Anatomy of a Neural Network (Simplified Explanation with Examples)

A **neural network** is like a team of problem solvers working together to understand and predict things based on the data they receive. It mimics the human brain, where neurons (small processing units) work together to process information.

Let us break it down step by step.

1. Layers of a Neural Network

A neural network has three main types of layers:

1.1 Input Layer

- This layer is like the **eyes and ears** of the network.
- It takes in the raw data (numbers, images, text, etc.) and passes it to the next layer.
- Each neuron in the input layer represents one feature of the data.

◆ Example:

If we are building a neural network to predict house prices, the input layer might have neurons representing:

- Size of the house (sq. ft.)
 - Number of bedrooms
 - Location rating
 - Age of the house
-

1.2 Hidden Layers

- These layers process the data by learning patterns and relationships.
- The more hidden layers a network has, the more complex patterns it can learn.
- Each neuron in a hidden layer takes input, applies a mathematical transformation, and passes it to the next layer.

◆ Example:

For house price prediction, hidden layers might learn:

- How house size affects price.
- How location influences value.
- How different features interact.

◆ Real-Life Example:

Think of hidden layers like the **chefs in a restaurant**.

- The **input layer** is the waiter taking orders.
 - The **hidden layers (chefs)** prepare the meal based on the order details.
 - The **output layer** delivers the final dish.
-

1.3 Output Layer

- This layer provides the final prediction or result.
- The number of neurons in this layer depends on the type of problem.

Types of Output Layers:

1. Binary Classification (Yes/No, True/False):

- Example: Spam detection (Spam or Not Spam)
- Uses **one neuron** with a **sigmoid activation function** (output is between 0 and 1).

2. Multiclass Classification (Multiple Categories):

- Example: Recognizing different types of animals (Cat, Dog, Horse).
- Uses **multiple neurons**, one for each category, with a **softmax activation function**.

3. Regression (Numerical Prediction):

- Example: Predicting house prices.
- Uses **one neuron** with a **linear activation function** (output is a real number).

◆ Real-Life Example:

The output layer is like the **cashier at a restaurant** who gives the final bill based on the order and what the chefs prepared.

2. Neurons (Processing Units)

- Each neuron receives input, processes it, and passes it forward.

- Neurons are connected by **weights**, which determine the importance of each input.

◆ **Example:**

In a spam detection email classifier:

- A neuron might check if the subject line contains the word "FREE".
 - Another neuron might check if the email contains too many links.
 - If both neurons detect spam-like features, they pass a strong signal forward.
-

3. Weights and Biases (How the Network Learns)

- **Weights:** Control the strength of connections between neurons.
- **Bias:** Helps adjust the output, making the network more flexible.

◆ **Example:**

- If a house's **location** is very important in determining its price, the network will assign a **high weight** to that input.
- If a house's **paint color** is not important, the network will assign a **low weight** to that input.

◆ **Real-Life Example:**

Think of weights like **teacher grading criteria**:

- Homework (20%)
 - Tests (50%)
 - Class participation (30%)
- Each factor is **weighted** differently in the final grade calculation.
-

4. Activation Functions (Decision Makers in the Network)

- Activation functions decide whether a neuron should be "activated" (send information forward) or not.
- Without activation functions, the network would just be performing basic linear calculations.

◆ **Types of Activation Functions:**

1. **Sigmoid:**

- Used for binary classification (Yes/No).
- Example: Detecting if an email is spam or not.

2. **ReLU (Rectified Linear Unit):**

- Used in hidden layers because it helps train deep networks efficiently.
- Example: Image recognition (detecting edges, colors, shapes).

3. Softmax:

- Used for multi-class classification problems.
- Example: Identifying different types of animals (dog, cat, bird).

◆ Real-Life Example:

Think of activation functions like **light switches** in a house:

- Some switches turn on a little (dim light – sigmoid).
 - Some turn on fully or not at all (on/off – ReLU).
-

5. Loss Function (How Wrong is the Prediction?)

- The loss function tells the neural network how far off its prediction is from the actual value.
- The goal is to minimize this error during training.

◆ Types of Loss Functions:

- **Binary Crossentropy:** Used for Yes/No classification.
- **Categorical Crossentropy:** Used for multi-class classification.
- **Mean Squared Error (MSE):** Used for predicting numbers (e.g., house prices).

◆ Real-Life Example:

- A loss function is like **Google Maps estimating your travel time**.
 - If it predicts you will take 30 minutes but you actually take 40, the loss function tells it to improve future estimates.
-

6. Optimizers (How the Network Learns from Mistakes)

- Optimizers adjust the weights and biases based on the loss function to improve predictions.

◆ Common Optimizers:

1. **Stochastic Gradient Descent (SGD):** Updates weights gradually.
2. **Adam Optimizer:** Adapts learning speed dynamically (most popular).

◆ Real-Life Example:

- Think of an optimizer as a **coach helping an athlete improve**.

- If a runner is too slow, the coach suggests changes in training.
 - The optimizer tweaks the neural network similarly.
-

7. Backpropagation (Learning from Mistakes)

- The network calculates how wrong it was, then updates weights to improve.
- This is done using **gradient descent**.

◆ Real-Life Example:

- Imagine a student taking a math test.
 - If they get a question wrong, the teacher shows them their mistake.
 - Next time, they try not to repeat the mistake.
-

Final Real-Life Example: Self-Driving Cars

A self-driving car uses a neural network to make decisions:

1. **Input Layer:** Sensors collect data (speed, road signs, objects).
2. **Hidden Layers:** Process patterns (detect red lights, pedestrians).
3. **Output Layer:** Decides whether to stop, slow down, or turn.
4. **Weights & Biases:** Adjust importance (pedestrians > lane markings).
5. **Loss Function:** If the car makes a bad decision, it learns from it.
6. **Optimizer:** Adjusts to improve future driving performance.

Introduction to Keras: Understanding Keras in AI & Deep Learning

1. What is Keras?

Keras is an **open-source deep learning framework** that provides a simple and intuitive interface for building neural networks. It is written in **Python** and is widely used in **Artificial Intelligence (AI) and Machine Learning (ML)** applications.

1.1 Why Use Keras?

- ✓ **User-Friendly** – Simple and easy to use for beginners.
- ✓ **Modular & Flexible** – Provides a high-level API for rapid development.
- ✓ **Runs on Multiple Backends** – Supports **TensorFlow, Theano, and Microsoft CNTK**.
- ✓ **Supports CPUs & GPUs** – Efficient computation for deep learning.
- ✓ **Pretrained Models** – Includes VGG, ResNet, MobileNet, etc., for transfer learning.

◆ Real-Life Analogy:

Keras is like a **modern kitchen with pre-built tools** that help you cook faster, while TensorFlow is like a **professional kitchen** where you have to do everything manually.

2. Why is Keras Important?

Keras **simplifies deep learning** by providing an easy-to-use interface while maintaining the power of advanced frameworks like **TensorFlow**.

◆ Key Benefits:

1. **Rapid Prototyping** – Build and test models quickly.
 2. **Scalability** – Works on small and large datasets.
 3. **Community Support** – Large user base and active development.
 4. **Integration with TensorFlow** – Keras is now part of TensorFlow as **tf.keras**.
-

3. How Keras Works?

Keras acts as a **high-level wrapper** that runs on top of deep learning frameworks such as:

Backend	Description
TensorFlow	Google's powerful deep learning framework (Most popular backend)
Theano	Early framework for deep learning (Now deprecated)
Microsoft CNTK	Deep learning framework by Microsoft (Less popular)

--> **Keras now primarily uses TensorFlow as its backend.**

4. Setting Up Keras

Before using Keras, we need to install it along with TensorFlow.

4.1 Installation

```
pip install tensorflow keras
```

4.2 Importing Keras

```
import keras
```

```
import tensorflow as tf
```

```
print(keras.__version__) # Check Keras version
```

```
print(tf.__version__) # Check TensorFlow version
```

5. Building a Neural Network with Keras

Keras provides an easy way to build **deep learning models** using a **Sequential API**.

5.1 Step-by-Step Guide to Building a Simple Neural Network

We will create a **basic neural network** to classify **handwritten digits (MNIST dataset)**.

Step 1: Import Libraries

```
from keras.models import Sequential
```

```
from keras.layers import Dense, Flatten
```

```
from keras.datasets import mnist
```

```
from keras.utils import to_categorical
```

Step 2: Load & Preprocess Data

```
# Load the MNIST dataset
```

```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

```
# Normalize images (scale pixel values between 0 and 1)
```

```
train_images = train_images / 255.0
```

```
test_images = test_images / 255.0
```

```
# Convert labels to categorical format (one-hot encoding)
```

```
train_labels = to_categorical(train_labels, 10)
```

```
test_labels = to_categorical(test_labels, 10)
```

Step 3: Create the Model

```
model = Sequential([
    Flatten(input_shape=(28, 28)), # Convert 2D images to 1D
    Dense(128, activation='relu'), # Hidden layer with 128 neurons
    Dense(10, activation='softmax') # Output layer with 10 neurons (for digits 0-9)
])
```

Step 4: Compile the Model

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

Step 5: Train the Model

```
model.fit(train_images, train_labels, epochs=5, batch_size=32, validation_split=0.2)
```

Step 6: Evaluate the Model

```
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test Accuracy: {test_acc}')
```

6. Keras Model APIs

Keras provides different ways to build models:

6.1 Sequential API (Easy & Quick)

- Builds models layer-by-layer in a **linear stack**.
- **Example:**

```
model = Sequential([
    Dense(64, activation='relu', input_shape=(784,)),
    Dense(10, activation='softmax')
])
```

6.2 Functional API (More Flexible)

- Allows creating complex architectures (multiple inputs/outputs).
- **Example:**

```
from keras.models import Model
```

```
from keras.layers import Input, Dense
```

```
input_layer = Input(shape=(784,))
```

```
hidden_layer = Dense(128, activation='relu')(input_layer)
```

```
output_layer = Dense(10, activation='softmax')(hidden_layer)
```

```
model = Model(inputs=input_layer, outputs=output_layer)
```

7. Key Components of Keras

Component	Description	Example
Layers	Building blocks of neural networks	Dense(128, activation='relu')
Optimizers	Adjust weights to minimize loss	optimizer='adam'
Loss Functions	Measure prediction error	loss='categorical_crossentropy'
Metrics	Evaluate model performance	metrics=['accuracy']

8. Pretrained Models in Keras (Transfer Learning)

Keras provides **pretrained models** that can be used for **image classification, object detection, and feature extraction**.

8.1 Using Pretrained Model (MobileNetV2)


```
from keras.applications import MobileNetV2
```

```
# Load MobileNetV2 model with pretrained weights
```

```
model = MobileNetV2(weights='imagenet')
```

✓ Used for Image Classification & Transfer Learning.

9. Applications of Keras

 Keras is widely used in various AI applications:

- ✓ **Image Recognition (Face Detection, Object Recognition)** – Used in **Self-driving Cars, Medical Imaging.**
- ✓ **Natural Language Processing (NLP)** – Used in **Chatbots, Speech Recognition.**
- ✓ **Recommendation Systems** – Used by **Netflix, Amazon, YouTube.**
- ✓ **Healthcare & Drug Discovery** – Used in **AI-based disease detection.**
- ✓ **Finance & Fraud Detection** – Used in **predicting stock prices & fraud detection.**

Here are some key topics:

- 1 **Keras Layers** – Understanding Dense, Convolutional, Recurrent layers, etc.
- 2 **Keras Optimizers** – Adam, SGD, RMSprop, and how they affect training.
- 3 **Keras Loss Functions** – How to choose the right loss function for classification & regression.
- 4 **Keras Callbacks** – How to use EarlyStopping, ModelCheckpoint, ReduceLROnPlateau, etc.
- 5 **Keras Functional API** – Building complex models with multiple inputs/outputs.
- 6 **Keras Transfer Learning** – Using pretrained models like VGG, ResNet, MobileNet.
- 7 **Keras for NLP** – Text processing, embedding layers, and LSTMs.
- 8 **Keras for Computer Vision** – Convolutional Neural Networks (CNNs).

Great! Let's go through all the important **Keras features** in detail, step by step.

1. Keras Layers

Layers are the **building blocks** of a neural network in Keras.

1.1 Types of Layers in Keras

Layer Type	Purpose	Example Usage
Dense (Fully Connected)	Basic neural network layer	Dense(128, activation='relu')
Conv2D (Convolutional Layer)	Extracts features from images	Conv2D(32, kernel_size=3, activation='relu')
MaxPooling2D (Pooling Layer)	Reduces spatial dimensions in CNNs	MaxPooling2D(pool_size=2)
Flatten	Converts multi-dimensional input to 1D	Flatten()
Dropout	Prevents overfitting	Dropout(0.5)

Layer Type	Purpose	Example Usage
LSTM/GRU (Recurrent Layer)	Used in sequence models (NLP, Time Series)	LSTM(64)

1.2 Example: Using Different Layers in Keras

```
from keras.models import Sequential
```

```
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout
```

```
# Define model
```

```
model = Sequential([
```

```
    Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=(28,28,1)),
```

```
    MaxPooling2D(pool_size=(2,2)),
```

```
    Flatten(),
```

```
    Dense(128, activation='relu'),
```

```
    Dropout(0.5),
```

```
    Dense(10, activation='softmax')
```

```
])
```

```
# Compile model
```

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# Model Summary
```

```
model.summary()
```

✓ Used in Image Classification, NLP, and other Deep Learning applications.

2. Keras Optimizers

Optimizers **adjust the weights** of the neural network to minimize the loss function.

2.1 Common Optimizers in Keras

Optimizer	Description	Best Use Case
SGD (Stochastic Gradient Descent)	Simple, slow but stable	Small datasets, linear regression
Adam (Adaptive Moment Estimation)	Most widely used, fast convergence	Works well on most problems
RMSprop	Designed for RNNs, adapts learning rates	Time-series, speech recognition
Adagrad	Adjusts learning rates dynamically	Sparse data, NLP tasks

2.2 Example: Using Different Optimizers

```
from keras.optimizers import Adam, SGD, RMSprop
```

```
# Compile model with different optimizers
```

```
model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy',  
metrics=['accuracy'])
```

✓ Adam is recommended for most deep learning applications.

3. Keras Loss Functions

Loss functions measure **how well the model is performing**.

3.1 Choosing the Right Loss Function

Problem Type	Loss Function
Binary Classification (Yes/No)	binary_crossentropy
Multiclass Classification	categorical_crossentropy
Regression (Predicting Numbers)	mean_squared_error

3.2 Example: Choosing a Loss Function

Binary Classification

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

Multiclass Classification

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

Regression

```
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
```

✓ **Choosing the right loss function improves model accuracy.**

4. Keras Callbacks

Callbacks are used to **monitor and control the training process**.

4.1 Common Callbacks

Callback	Purpose
EarlyStopping	Stops training if the model stops improving
ModelCheckpoint	Saves the best model during training
ReduceLROnPlateau	Reduces learning rate when the model stops improving

4.2 Example: Using Callbacks

```
from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau
```

```
callbacks = [
```

```
    EarlyStopping(monitor='val_loss', patience=5),
```

```
    ModelCheckpoint('best_model.h5', save_best_only=True),
```

```
    ReduceLROnPlateau(monitor='val_loss', factor=0.1, patience=3)
```

```
]
```

```
model.fit(train_images, train_labels, epochs=50, batch_size=32, validation_split=0.2,
callbacks=callbacks)
```

✓ **Callbacks help prevent overfitting and speed up training.**

5. Keras Functional API

The **Functional API** allows building complex models with **multiple inputs/outputs**.

5.1 Example: Functional API Model

```
from keras.models import Model

from keras.layers import Input, Dense

# Input Layer
input_layer = Input(shape=(784,))

# Hidden Layers
hidden_layer = Dense(128, activation='relu')(input_layer)

# Output Layer
output_layer = Dense(10, activation='softmax')(hidden_layer)

# Create Model
model = Model(inputs=input_layer, outputs=output_layer)

# Compile Model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.summary()
```

✓ Used in advanced deep learning applications like multi-task learning.

6. Keras Transfer Learning

Transfer learning uses **pretrained models** to improve performance on new datasets.

6.1 Using a Pretrained Model (MobileNetV2)

```
from keras.applications import MobileNetV2
from keras.models import Model
from keras.layers import Dense, Flatten

# Load MobileNetV2 model
base_model = MobileNetV2(weights='imagenet', include_top=False,
input_shape=(224,224,3))

# Freeze the base model layers
for layer in base_model.layers:
    layer.trainable = False

# Add new layers
x = Flatten()(base_model.output)
x = Dense(128, activation='relu')(x)
output_layer = Dense(10, activation='softmax')(x)

# Create model
model = Model(inputs=base_model.input, outputs=output_layer)

# Compile model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

model.summary()
```

✓ Used for Image Recognition & Object Detection.

7. Keras for NLP (Text Processing)

Keras provides tools for **Natural Language Processing (NLP)**.

7.1 Example: Using Embedding Layer for NLP

```
from keras.layers import Embedding, LSTM

model = Sequential([
    Embedding(input_dim=10000, output_dim=128, input_length=100),
    LSTM(64),
    Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

✓ Used in Chatbots, Sentiment Analysis, and Speech Recognition.

8. Keras for Computer Vision (CNNs)

CNNs are used for **image classification and object detection**.

8.1 Example: CNN Model

```
model = Sequential([
    Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=(64,64,3)),
    MaxPooling2D(pool_size=(2,2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

✓ Used in Face Recognition, Medical Imaging, and Self-Driving Cars.

- 1 **Keras Sequential API** – How to build deep learning models step by step.
- 2 **Keras Functional API** – Creating complex models with multiple inputs/outputs.
- 3 **Keras CNNs (Computer Vision)** – Building and training a Convolutional Neural Network (CNN).
- 4 **Keras NLP (Text Processing & LSTMs)** – Using Keras for text classification, sentiment analysis, and RNNs.
- 5 **Keras Transfer Learning** – Using pretrained models like VGG16, ResNet, and MobileNet.
- 6 **Keras Hyperparameter Tuning** – Optimizing model performance with Keras Tuner.
- 7 **Keras Callbacks & Model Saving** – Using EarlyStopping, ModelCheckpoint, and saving/loading models.

Awesome! Let's go through **all key Keras topics** in detail with **step-by-step tutorials** and **Python examples**. 🚀

1 Keras Sequential API

The **Sequential API** is the easiest way to build deep learning models in Keras. It allows stacking layers in a **linear fashion** (one after another).

1.1 Example: Build a Simple Neural Network (Sequential API)

```
from keras.models import Sequential
from keras.layers import Dense

# Create a Sequential model
model = Sequential([
    Dense(128, activation='relu', input_shape=(784,)), # Hidden layer with 128 neurons
    Dense(10, activation='softmax') # Output layer for 10 classes
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Summary of the model
model.summary()

✓ Best for simple feedforward networks.
```

2 Keras Functional API

The **Functional API** allows creating **complex architectures**, like multiple inputs/outputs and shared layers.

2.1 Example: Multi-Input Model (Functional API)

```
from keras.models import Model

from keras.layers import Input, Dense, concatenate

# Input layers
input_a = Input(shape=(64,))
input_b = Input(shape=(32,))

# Hidden layers
hidden_a = Dense(32, activation='relu')(input_a)
hidden_b = Dense(32, activation='relu')(input_b)

# Merge both branches
merged = concatenate([hidden_a, hidden_b])
output = Dense(1, activation='sigmoid')(merged)

# Create model
model = Model(inputs=[input_a, input_b], outputs=output)

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Summary of the model
model.summary()

✓ Best for models with multiple inputs and outputs.
```

3 Keras CNNs (Computer Vision)

Convolutional Neural Networks (CNNs) are used for **image classification, object detection, and facial recognition.**

3.1 Example: CNN for Image Classification

```
from keras.models import Sequential

from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

# Define CNN model
model = Sequential([
    Conv2D(32, kernel_size=(3,3), activation='relu', input_shape=(64,64,3)),
    MaxPooling2D(pool_size=(2,2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Model summary
model.summary()

✓ Used in Face Recognition, Object Detection, and Medical Imaging.
```

4 Keras NLP (Text Processing & LSTMs)

Keras provides tools for **Natural Language Processing (NLP)** such as **word embeddings** and **LSTMs**.

4.1 Example: Using Embedding Layer for Text Classification

```
from keras.models import Sequential

from keras.layers import Embedding, LSTM, Dense

# Define NLP model
model = Sequential([
    Embedding(input_dim=10000, output_dim=128, input_length=100),
    LSTM(64),
    Dense(1, activation='sigmoid')
])

# Compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Model summary
model.summary()

✓ Used in Chatbots, Sentiment Analysis, and Speech Recognition.
```

5 Keras Transfer Learning

Transfer learning allows **using pretrained models** like **VGG16, ResNet, MobileNet** for image classification.

5.1 Example: Using MobileNetV2 for Transfer Learning

```
from keras.applications import MobileNetV2

from keras.models import Model

from keras.layers import Flatten, Dense

# Load pretrained MobileNetV2 model
```

```
base_model = MobileNetV2(weights='imagenet', include_top=False,
input_shape=(224,224,3))

# Freeze the base model layers
for layer in base_model.layers:
    layer.trainable = False

# Add new layers
x = Flatten()(base_model.output)
x = Dense(128, activation='relu')(x)
output_layer = Dense(10, activation='softmax')(x)

# Create model
model = Model(inputs=base_model.input, outputs=output_layer)

# Compile model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Model summary
model.summary()
```

✓ Used in Image Recognition, Object Detection, and Medical Diagnosis AI.

6 Keras Hyperparameter Tuning

Tuning hyperparameters like **learning rate, batch size, and number of layers** improves model performance.

6.1 Example: Hyperparameter Tuning with Keras Tuner

```
import keras_tuner as kt
from keras.models import Sequential
from keras.layers import Dense
```

```
# Function to build the model

def build_model(hp):
    model = Sequential()

    model.add(Dense(hp.Int('units', min_value=32, max_value=256, step=32),
activation='relu', input_shape=(784,)))

    model.add(Dense(10, activation='softmax'))

    model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

    return model


# Define tuner

tuner = kt.RandomSearch(build_model, objective='val_accuracy', max_trials=5)


# Run the search

tuner.search(train_images, train_labels, epochs=10, validation_split=0.2)

✓ Used to find the best hyperparameters automatically.
```

7 Keras Callbacks & Model Saving

Callbacks help **control training** and **save the best model**.

7.1 Example: Using Callbacks

python

CopyEdit

```
from keras.callbacks import EarlyStopping, ModelCheckpoint
```

```
callbacks = [
    EarlyStopping(monitor='val_loss', patience=5),
    ModelCheckpoint('best_model.h5', save_best_only=True)
]
```

Train model with callbacks

```
model.fit(train_images, train_labels, epochs=50, batch_size=32, validation_split=0.2,
callbacks=callbacks)
```

✓ Used to prevent overfitting and save the best model.

7.2 Saving and Loading a Model

Save model

```
model.save('my_model.h5')
```

Load model

```
from keras.models import load_model
```

```
loaded_model = load_model('my_model.h5')
```

Check model summary

```
loaded_model.summary()
```

✓ Ensures models can be reused without retraining.

8 Full Example: Building & Training a Keras Model

Let's build and train a complete deep learning model using **Keras Sequential API**.

```
from keras.models import Sequential
```

```
from keras.layers import Dense, Flatten
```

```
from keras.datasets import mnist
```

```
from keras.utils import to_categorical
```

Load the dataset

```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

Normalize images

```
train_images, test_images = train_images / 255.0, test_images / 255.0
```

Convert labels to categorical format

```
train_labels = to_categorical(train_labels, 10)
```

```
test_labels = to_categorical(test_labels, 10)
```


```
# Define model
model = Sequential([
    Flatten(input_shape=(28, 28)),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

# Compile model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train model
model.fit(train_images, train_labels, epochs=5, batch_size=32, validation_split=0.2)

# Evaluate model
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test Accuracy: {test_acc}')
```

✓ Used for digit classification in the MNIST dataset.

 **Keras provides a wide range of functions across different modules. Below is a categorized list of the most used functions in the Keras library.**

1. Model Creation & Architecture

- `tf.keras.Sequential()`
 - `tf.keras.Model()`
 - `tf.keras.models.load_model()`
 - `tf.keras.models.clone_model()`
 - `tf.keras.models.model_from_json()`
 - `tf.keras.models.model_from_yaml()`
-

2. Layers

- `tf.keras.layers.Dense()`
 - `tf.keras.layers.Conv2D()`
 - `tf.keras.layers.MaxPooling2D()`
 - `tf.keras.layers.Flatten()`
 - `tf.keras.layers.Dropout()`
 - `tf.keras.layers.BatchNormalization()`
 - `tf.keras.layers.LSTM()`
 - `tf.keras.layers.GRU()`
 - `tf.keras.layers.Embedding()`
 - `tf.keras.layers.Input()`
 - `tf.keras.layers.Reshape()`
 - `tf.keras.layers.ReLU()`
 - `tf.keras.layers.Softmax()`
 - `tf.keras.layers.Activation()`
-

3. Model Compilation & Training

- `tf.keras.Model.compile()`
 - `tf.keras.Model.fit()`
 - `tf.keras.Model.fit_generator()`
 - `tf.keras.Model.evaluate()`
 - `tf.keras.Model.predict()`
 - `tf.keras.Model.train_on_batch()`
 - `tf.keras.Model.test_on_batch()`
 - `tf.keras.Model.predict_on_batch()`
-

4. Optimizers

- `tf.keras.optimizers.Adam()`
- `tf.keras.optimizers.SGD()`

- `tf.keras.optimizers.RMSprop()`
 - `tf.keras.optimizers.Adagrad()`
 - `tf.keras.optimizers.Adamax()`
 - `tf.keras.optimizers.Nadam()`
-

5. Loss Functions

- `tf.keras.losses.MeanSquaredError()`
 - `tf.keras.losses.BinaryCrossentropy()`
 - `tf.keras.losses.CategoricalCrossentropy()`
 - `tf.keras.losses.SparseCategoricalCrossentropy()`
 - `tf.keras.losses.Hinge()`
 - `tf.keras.losses.KLDivergence()`
-

6. Metrics

- `tf.keras.metrics.Accuracy()`
 - `tf.keras.metrics.Precision()`
 - `tf.keras.metrics.Recall()`
 - `tf.keras.metrics.AUC()`
 - `tf.keras.metrics.MeanSquaredError()`
 - `tf.keras.metrics.CategoricalAccuracy()`
 - `tf.keras.metrics.SparseCategoricalAccuracy()`
-

7. Callbacks

- `tf.keras.callbacks.EarlyStopping()`
 - `tf.keras.callbacks.ModelCheckpoint()`
 - `tf.keras.callbacks.ReduceLROnPlateau()`
 - `tf.keras.callbacks.TensorBoard()`
 - `tf.keras.callbacks.LearningRateScheduler()`
 - `tf.keras.callbacks.CSVLogger()`
-

8. Preprocessing & Data Augmentation

Image Processing

- `tf.keras.preprocessing.image.load_img()`
- `tf.keras.preprocessing.image.img_to_array()`
- `tf.keras.preprocessing.image.array_to_img()`
- `tf.keras.preprocessing.image.ImageDataGenerator()`

Text Processing

- `tf.keras.preprocessing.text.Tokenizer()`
 - `tf.keras.preprocessing.sequence.pad_sequences()`
-

9. Regularization

- `tf.keras.regularizers.l1()`
 - `tf.keras.regularizers.l2()`
 - `tf.keras.regularizers.l1_l2()`
-

10. Saving & Loading

- `tf.keras.models.save_model()`
- `tf.keras.models.load_model()`
- `tf.keras.models.model_from_json()`
- `tf.keras.models.model_from_yaml()`

Detailed Explanation of Keras Functions

Keras provides a wide range of functions across different modules. Below is a detailed explanation of each function, its use case, and when to use it.

1. Model Creation & Architecture

`tf.keras.Sequential()`

- **Use:** Creates a linear stack of layers.
- **Used in:** Simple feedforward neural networks.

Example:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
model = Sequential([
    Dense(64, activation='relu', input_shape=(32,)),
    Dense(10, activation='softmax')
])
```

tf.keras.Model()

- **Use:** Defines a model using the Functional API.
- **Used in:** Complex models with multiple inputs/outputs.
- **Example:**

```
from tensorflow.keras.layers import Input, Dense
from tensorflow.keras.models import Model
inputs = Input(shape=(32,))
x = Dense(64, activation="relu")(inputs)
outputs = Dense(10, activation="softmax")(x)
model = Model(inputs, outputs)
```

tf.keras.models.load_model()

- **Use:** Loads a previously saved model.
- **Used in:** Deploying trained models.

Example:

```
from tensorflow.keras.models import load_model
model = load_model('my_model.h5')
```

tf.keras.models.clone_model()

- **Use:** Creates an exact copy of a model architecture.
- **Used in:** Transfer learning without copying weights.

Example:

```
from tensorflow.keras.models import clone_model
model_copy = clone_model(model)
```

tf.keras.models.model_from_json()

- **Use:** Loads a model from a JSON file.
- **Used in:** When sharing model architecture.

Example:

```
json_string = model.to_json()
new_model = tf.keras.models.model_from_json(json_string)
```

tf.keras.models.model_from_yaml()

- Deprecated in newer versions of Keras.

2. Layers

tf.keras.layers.Dense()

- **Use:** Fully connected layer.
- **Used in:** Feedforward neural networks.

tf.keras.layers.Conv2D()

- **Use:** Applies convolutional filters.
- **Used in:** Convolutional Neural Networks (CNNs).

tf.keras.layers.MaxPooling2D()

- **Use:** Reduces spatial size.
- **Used in:** CNNs for downsampling.

tf.keras.layers.Flatten()

- **Use:** Converts multi-dimensional input into 1D.
- **Used in:** Connecting CNNs to fully connected layers.

tf.keras.layers.Dropout()

- **Use:** Prevents overfitting.
- **Used in:** Any deep learning model.

tf.keras.layers.BatchNormalization()

- **Use:** Normalizes activations.
- **Used in:** Speeding up training.

tf.keras.layers.LSTM()

- **Use:** Long Short-Term Memory unit.
- **Used in:** Time series, NLP.

tf.keras.layers.GRU()

- **Use:** Gated Recurrent Unit.
- **Used in:** Faster alternative to LSTM.

tf.keras.layers.Embedding()

- **Use:** Converts words to vector representations.
- **Used in:** NLP models.

tf.keras.layers.Input()

- **Use:** Defines an input layer.

tf.keras.layers.Reshape()

- **Use:** Reshapes tensor without changing data.

tf.keras.layers.ReLU()

- **Use:** Applies Rectified Linear Activation function.

tf.keras.layers.Softmax()

- **Use:** Converts logits to probabilities.

tf.keras.layers.Activation()

- **Use:** Applies any activation function.

3. Model Compilation & Training

tf.keras.Model.compile()

- **Use:** Configures the model for training.
- **Used in:** Every model before training.

tf.keras.Model.fit()

- **Use:** Trains the model.
- **Used in:** Supervised learning models.

tf.keras.Model.evaluate()

- **Use:** Tests model performance.

tf.keras.Model.predict()

- **Use:** Makes predictions.

4. Optimizers

tf.keras.optimizers.Adam()

- **Use:** Adaptive learning rate optimizer.
- **Used in:** Most deep learning models.

tf.keras.optimizers.SGD()

- **Use:** Stochastic Gradient Descent.
- **Used in:** Classic ML models.

tf.keras.optimizers.RMSprop()

- **Use:** Adaptive learning for RNNs.

5. Loss Functions

tf.keras.losses.MeanSquaredError()

- **Use:** Measures squared loss.
- **Used in:** Regression problems.

tf.keras.losses.BinaryCrossentropy()

- **Use:** Binary classification loss.

tf.keras.losses.CategoricalCrossentropy()

- **Use:** Multi-class classification loss.

6. Metrics

tf.keras.metrics.Accuracy()

- **Use:** Measures accuracy.

tf.keras.metrics.Precision()

- **Use:** Measures precision.

7. Callbacks

`tf.keras.callbacks.EarlyStopping()`

- **Use:** Stops training when performance degrades.

`tf.keras.callbacks.ModelCheckpoint()`

- **Use:** Saves model checkpoints.

8. Preprocessing & Data Augmentation

`tf.keras.preprocessing.image.load_img()`

- **Use:** Loads an image from disk.

`tf.keras.preprocessing.text.Tokenizer()`

- **Use:** Tokenizes text.

9. Regularization

`tf.keras.regularizers.l1()`

- **Use:** L1 regularization.

`tf.keras.regularizers.l2()`

- **Use:** L2 regularization.

10. Saving & Loading

`tf.keras.models.save_model()`

- **Use:** Saves a trained model.

`tf.keras.models.load_model()`

- **Use:** Loads a trained model.

tf.keras.layers.Dense(): A Detailed Explanation

Overview

tf.keras.layers.Dense() is a fully connected (or dense) layer in a neural network, where each neuron in the layer receives input from all neurons in the previous layer. It is the fundamental building block of feedforward neural networks.

Syntax

```
1  tf.keras.layers.Dense(  
2      units,  
3      activation=None,  
4      use_bias=True,  
5      kernel_initializer='glorot_uniform',  
6      bias_initializer='zeros',  
7      kernel_regularizer=None,  
8      bias_regularizer=None,  
9      activity_regularizer=None,  
10     kernel_constraint=None,  
11     bias_constraint=None  
12 )
```

Parameters

- units (int): The number of neurons (or output dimensions) in the layer.
- activation (str or function): Activation function to apply (e.g., 'relu', 'sigmoid', 'softmax'). Default is None (linear activation).
- use_bias (bool): Whether the layer uses a bias vector.
- kernel_initializer (str or function): Initialization method for weights (default: 'glorot_uniform').
- bias_initializer (str or function): Initialization method for bias (default: 'zeros').
- kernel_regularizer (function): Regularization method for weights (e.g., tf.keras.regularizers.l2(0.01)).
- bias_regularizer (function): Regularization method for bias.
- activity_regularizer (function): Regularization method for outputs.
- kernel_constraint (function): Constraint on weights (e.g., tf.keras.constraints.max_norm(2.0)).
- bias_constraint (function): Constraint on bias.

Use Cases of `tf.keras.layers.Dense()`

1. Basic Feedforward Neural Network

A Dense layer is widely used in fully connected neural networks.

Example: Simple Classification Model

```
1 import tensorflow as tf
2 from tensorflow.keras.models import Sequential
3 from tensorflow.keras.layers import Dense
4
5 # Define a simple feedforward network
6 model = Sequential([
7     Dense(64, activation='relu', input_shape=(10,)), # Hidden layer with 64 neurons
8     Dense(32, activation='relu'), # Another hidden layer
9     Dense(1, activation='sigmoid') # Output layer for binary classification
10 ])
11
12 # Compile the model
13 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
14
15 # Summary of the model
16 model.summary()
17
```

Use Case:

- Used in binary classification problems such as spam detection, fraud detection, and medical diagnosis.

2. Multi-Class Classification

For multi-class classification, the output layer uses softmax activation.

Example: Multi-Class Classification

```
1 model = Sequential([
2     Dense(128, activation='relu', input_shape=(20,)), # Hidden layer
3     Dense(64, activation='relu'), # Another hidden layer
4     Dense(10, activation='softmax') # Output layer (10 classes)
5 ])
6
7 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
8
```

Use Case:

- Used in image classification, sentiment analysis, and handwritten digit recognition (MNIST).

3. Regression Model

For regression problems, Dense layers with **no activation** in the output layer are used.

Example: Regression Model

```
1 model = Sequential([
2     Dense(64, activation='relu', input_shape=(15,)), # Hidden layer
3     Dense(32, activation='relu'), # Another hidden layer
4     Dense(1) # Output layer (No activation for regression)
5 ])
6
7 model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
```

Use Case:

- Used in predicting house prices, stock market predictions, and sales forecasting.
-

4. Autoencoder

Dense layers are used in **autoencoders** for dimensionality reduction.

Example: Autoencoder

```
1 input_dim = 100
2 encoding_dim = 32
3
4 # Encoder
5 encoder = Sequential([
6     Dense(encoding_dim, activation='relu', input_shape=(input_dim,))
7 ])
8
9 # Decoder
10 decoder = Sequential([
11     Dense(input_dim, activation='sigmoid')
12 ])
13
14 autoencoder = Sequential([encoder, decoder])
15
16 autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

Use Case:

- Used in anomaly detection, image denoising, and feature extraction.
-

5. Transfer Learning (Fully Connected Layers on Pretrained Models)

Dense layers are often added to **pretrained convolutional networks** (like VGG16) for custom classification.

Example: Transfer Learning

```

1  from tensorflow.keras.applications import VGG16
2  from tensorflow.keras.models import Model
3
4  # Load pre-trained VGG16 model (without the classification head)
5  base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
6
7  # Flatten output and add Dense layers for custom classification
8  x = tf.keras.layers.Flatten()(base_model.output)
9  x = Dense(256, activation='relu')(x)
10 x = Dense(10, activation='softmax')(x) # Output layer for 10 classes
11
12 # Define new model
13 model = Model(inputs=base_model.input, outputs=x)
14
15 # Freeze base model weights
16 for layer in base_model.layers:
17     layer.trainable = False
18
19 model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
20 |

```

Use Case:

- Used in medical image analysis, face recognition, and object detection.

Best Practices When Using `tf.keras.layers.Dense()`

1. **Use ReLU Activation in Hidden Layers:** Helps avoid vanishing gradient problems.
 2. **Use Softmax for Multi-Class Output:** Converts logits into probabilities.
 3. **Use No Activation for Regression Outputs:** Ensures continuous output.
 4. **Apply Dropout and Batch Normalization:** Helps prevent overfitting.
 5. **Use Regularization (l1, l2):** Helps improve generalization.
 6. **Ensure Proper Input Shape:** The first Dense layer must define `input_shape`.
-

Summary

Scenario	Example Architecture	Activation
Binary Classification	[Dense(64, 'relu'), Dense(1, 'sigmoid')]	'sigmoid'
Multi-Class Classification	[Dense(128, 'relu'), Dense(10, 'softmax')]	'softmax'
Regression	[Dense(64, 'relu'), Dense(1)]	None
Autoencoder	[Dense(32, 'relu'), Dense(100, 'sigmoid')]	'relu', 'sigmoid'
Transfer Learning	VGG16 + Dense(256, 'relu') + Dense(10, 'softmax')	'relu', 'softmax'

Conclusion

- Dense() is a core layer in deep learning, commonly used in **classification, regression, autoencoders, and transfer learning**.
- Understanding activation functions, input shapes, and optimizers is **crucial** for its effective use.
- Regularization techniques help improve **generalization and prevent overfitting**.