

## Ex-10 Lab Manual: Implement a Recurrent Neural Network (RNN) for IMDB Movie Review Classification

---

### 1. Objective

To implement and train a **Recurrent Neural Network (RNN)** for classifying movie reviews as **positive or negative** using the **IMDB dataset**. The RNN model will process sequential text data and learn relationships between words over time.

### 2. Introduction to RNNs

A **Recurrent Neural Network (RNN)** is a type of neural network that is designed to handle **sequential data**, making it well-suited for tasks such as text classification, speech recognition, and time-series forecasting. Unlike traditional feedforward neural networks, RNNs **retain information from previous inputs**, allowing them to capture the context in text sequences.

#### Key Features of RNNs:

- Maintain memory of past inputs through **hidden states**.
- Process variable-length sequences efficiently.
- Suitable for NLP tasks like **text classification, machine translation, and sentiment analysis**.

In this lab, we will implement an **RNN using Long Short-Term Memory (LSTM) cells**, which help mitigate the vanishing gradient problem common in traditional RNNs.

---

### 3. System Requirements

#### Hardware Requirements:

- Computer with at least **4GB RAM** (8GB recommended)
- GPU support for faster training (optional but recommended)

#### Software Requirements:

- Python (>=3.6)
- TensorFlow/Keras
- NumPy, Matplotlib (for data processing and visualization)

Install the required libraries using:

```
pip install numpy tensorflow matplotlib
```

---

## 4. Step-by-Step Procedure

### Step 1: Import Required Libraries

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense
import matplotlib.pyplot as plt
```

### Step 2: Load and Preprocess IMDB Dataset

The IMDB dataset consists of **50,000 movie reviews**, labeled as **positive (1)** or **negative (0)**.

# Define parameters

```
vocab_size = 10000 # Number of unique words to consider
```

```
max_length = 100 # Maximum words per review
```

```
embedding_dim = 32 # Size of word embeddings
```

# Load IMDB dataset (only the top 10,000 words)

```
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)
```

# Pad sequences to ensure uniform length

```
x_train = pad_sequences(x_train, maxlen=max_length, padding='post')
```

```
x_test = pad_sequences(x_test, maxlen=max_length, padding='post')
```

### Step 3: Define the RNN Model with LSTM Layers

```
model = Sequential([
```

```
    Embedding(input_dim=vocab_size, output_dim=embedding_dim,
input_length=max_length),
```

```
    LSTM(64, return_sequences=False), # LSTM layer with 64 units
```

```
Dense(1, activation='sigmoid') # Binary classification output
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.summary()
```

#### Step 4: Train the Model

```
history = model.fit(x_train, y_train, epochs=5, batch_size=32, validation_data=(x_test,
y_test))
```

#### Step 5: Evaluate the Model

```
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test Accuracy: {test_acc:.4f}")
```

#### Step 6: Plot Training History

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

---

### 5. Observations and Results

- The **RNN with LSTM** successfully classifies IMDB movie reviews as positive or negative.
- Training accuracy should improve over epochs, while validation accuracy indicates generalization.
- The **LSTM layer** captures long-range dependencies, improving classification performance compared to traditional models.

---

### 6. Troubleshooting & Optimization

- **Overfitting?** Apply dropout layers to reduce overfitting.

- **Low accuracy?** Increase the LSTM units or embedding dimensions.
  - **Slow training?** Use a GPU for faster computation.
- 

## 7. Additional Tasks

- Experiment with different **RNN architectures**, such as **GRUs**.
  - Use **pre-trained word embeddings** like GloVe or Word2Vec.
  - Extend the model to **multi-class classification** with more sentiment categories.
- 

## 8. Conclusion

RNNs, especially **LSTMs**, are powerful for text classification tasks like sentiment analysis. They capture contextual meaning and dependencies between words, making them highly effective for NLP applications.

---

**End of Lab Manual**

# what is Recurrent Neural Network

## Recurrent Neral Network (RNN) - A Detailed Explanation

### 1. Introduction

A **Recurrent Neural Network (RNN)** is a type of artificial neural network designed to process **sequential data**, such as text, speech, and time series data. Unlike traditional **feedforward neural networks**, RNNs have a mechanism to retain memory of previous inputs, making them useful for tasks where **context and order** matter.

#### Example Use Cases:

- ✅ **Text classification** (e.g., sentiment analysis, spam detection)
  - ✅ **Speech recognition** (e.g., Siri, Google Assistant)
  - ✅ **Machine translation** (e.g., Google Translate)
  - ✅ **Time series forecasting** (e.g., stock price prediction)
- 

### 2. Why Use RNNs?

Traditional neural networks, like **fully connected (dense) networks**, treat each input independently, **ignoring sequence relationships**. RNNs solve this problem by introducing a **hidden state** that **remembers past inputs**, enabling the model to understand sequential dependencies.

#### Comparison of Neural Networks

Model Type	Processes Sequences?	Remembers Context?	Common Use Case
Feedforward NN	❌ No	❌ No	Image classification
Recurrent NN	✅ Yes	✅ Yes	Text & speech processing
LSTM / GRU	✅ Yes	✅ Better	Long-term dependencies

---

### 3. How RNNs Work

An RNN processes an **input sequence step-by-step**, maintaining a hidden state  $h_t$  that gets updated at each step based on the **current input** and the **previous hidden state**.

#### Mathematical Representation

At each time step  $t$ :

$$h_t = f(W_h \cdot h_{t-1} + W_x \cdot x_t + b)$$

Where:

- $h_t$  is the hidden state at time  $t$
  - $W_h$  and  $W_x$  are weight matrices
  - $x_t$  is the current input
  - $b$  is a bias term
  - $f$  is an activation function (usually **tanh** or **ReLU**)
- 

### 4. Types of RNN Architectures

#### 4.1 Basic RNN

- The simplest form of RNN, but suffers from the **vanishing gradient problem**, making it hard to capture long-term dependencies.

#### 4.2 Long Short-Term Memory (LSTM)

- A special type of RNN that solves the vanishing gradient problem using **gates** (forget, input, and output gates).
- Ideal for **long-range dependencies** in text and speech data.

#### 4.3 Gated Recurrent Unit (GRU)

- A simplified version of LSTM with fewer parameters, making it faster.
  - Works well for short and medium-length sequences.
-

## 5. Advantages & Limitations of RNNs

Feature	RNNs	LSTMs/GRUs
Handles sequences?	✓ Yes	✓ Yes
Short-term memory?	✓ Good	✓ Excellent
Long-term memory?	✗ Weak	✓ Strong
Training efficiency	✓ Fast	✗ Slower (more parameters)

### ● Limitations of Basic RNNs:

- **Vanishing Gradient Problem:** Difficulty in learning long-term dependencies.
  - **Exploding Gradients:** Large gradients can make the model unstable.
  - **Slow Training:** Sequential processing limits parallel computation.
- 

## 6. Implementing an RNN in Python (Using TensorFlow/Keras)

### Step 1: Import Dependencies

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN, Dense, Embedding
```

### Step 2: Define the RNN Model

```
model = Sequential([
    Embedding(input_dim=5000, output_dim=32, input_length=100), # Word
    embeddings
    SimpleRNN(64, return_sequences=False), # RNN Layer
    Dense(1, activation='sigmoid') # Output layer for binary classification
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
model.summary()
```